

1 HOMEWORK SUBMISSION FORMAT

The homework assignment will consist of two parts:

- **Written Assignment:** The part of the assignment will contain all of the short answer questions. You will submit a pdf of your responses to the Canvas assignment titled: Written Assignment.
- **Coding Assignment:** This part of the assignment will contain all codes. It will be uploaded to the auto-grader as a tarball like usual.

Remark 1: List your section number (i.e. 483 or 583) at the top of your PDF file.

Remark 2: When you tar your files only include all .cpp files, .hpp files and the makefile. Do not tar the grader candidate folder, or the entire grader! Your command for tarring files should look something like this

```
tar -czf ps5.tgz *.cpp *.hpp Makefile
```

2 Preliminaries

2.1 Bandwidth and Hierarchical Memory Parameters

One important basic characterization that you can make of your machine is to measure the bandwidth for various sizes of data and with various types of data movement instructions. Since many computations are bandwidth bound, this can give you information with which you can estimate performance bounds.

A docker image called `amath583/bandwidth` that will profile your computer and create a bandwidth plot has been put on docker hub. When running this image will put its results into a subdirectory called "BW". So, before running this image, first create a subdirectory named BW in the directory you normally work in with docker. Also, remember to run `docker pull <image name>` to get the latest version of the docker image. Now you can run the image. That is:

```
$ cd <your working directory>
$ mkdir BW
$ docker run -v <your working directory>:/home/amath583/work amath583/bandwidth
```

The `<your working directory>` stands for the location on your host machine where you have been doing your work under docker. The bandwidth program will run for quite some time – 10-20 minutes – during which time your computer should have as little load on it as possible (other than the bandwidth program itself).

When the program completes, there will be two files in the BW directory – `bandwidth.png` and `bandwidth.csv`. The file `bandwidth.png` is an image of a graph showing data transfer rates for various types of operations and various types of operations. The file `bandwidth.csv` is the raw data that is plotted in `bandwidth.png`, as comma separated values. You can use that data to do your own investigations about your computer.

In looking at the bandwidth graph you will see one or more places where the transfer rate drops suddenly. For small amounts of data, the data can fit in cache closer to the CPU – and hence data transfer rates will be higher. The data transfer rates will fall off fairly quickly once that data no longer completely fits into cache. Based on the bandwidth plot obtained for your computer, how many levels of cache does it have? How large are each of the caches (at what sizes does the bandwidth fall)? If possible, check your estimates against published information for the microprocessor that your system is using.

You may also notice that write operations tend to be slower than the corresponding read operations.

2.2 Roofline Parameters

A more refined characterization of your machine that accounts for compute power as well as bandwidth is the roofline model, which measures performance not simply a function of the amount of data being transferred, nor just raw compute power, but rather as a function of numerical intensity. As discussed in lecture, the roofline model depends on both the memory and CPU characteristics of the system being modeled. Although one could look up various parameters about your system, it is more accurate (and informative) to measure the actual system.

Two docker images that will profile your computer and create a roofline plot has been put on docker hub. One image is called `amath583/ert`, which will create a roofline plot based on a single CPU core of your machine. Another is called `amath583/ert.openmp`, which will create a roofline plot based on all CPU cores of your machine. Remember to run `docker pull <image name>` to get the latest version of the docker images. When running both images will put their results into a subdirectory called “ERT”. So, before running this image, first create a subdirectory named ERT in the directory you normally work in with docker. Then you can run both images. To run the `amath583/ert` image:

```
$ cd <your working directory>
$ mkdir ERT
$ docker run -v <your working directory>:/home/amath583/work amath583/ert
```

The `<your working directory>` stands for the location on your host machine where you have been doing your work under docker.

Note that as with bandwidth measurements, the roofline characterization will run for quite some time – 10-20 minutes – during which time your computer should have as little load on it as possible. I.e., this docker image should be the only thing running on your machine. The traditional activity when waiting for a long-running computational task is to get a cup of coffee.

When you are back from coffee and the job is completed, it will have created a small trove of the data it generated. Now you can run the `amath583/ert.openmp` image similarly:

```
$ cd <your working directory>
$ docker run -v <your working directory>:/home/amath583/work amath583/ert.openmp
```

The `<your working directory>` stands for the location on your host machine where you have been doing your work under docker. Note that you should let `amath583/ert` finish before you run `amath583/ert.openmp`.

Once both images finish, you will find results in the subdirectories named `Sequential` and `OpenMP`. In each of those subdirectories will be one or more subdirectories with the name `Run.001`. Directly in the `Run.001` subdirectories will be a file `roofline.pdf`, which contains a graph of the roofline model for your computer. **NB:** If you do multiple runs of the roofline profiler, delete the subdirectory `Run.001` each time so that the profiler has a fresh directory to write to and there won't be old or incomplete results to confuse you or to confuse the profiler.

3 Warm Up

3.1 Futures

C++ provides the `std::async()` mechanism to enable execution of a function call as an asynchronous task. Its semantics are important. Consider the example we have been using in lecture.

```
typedef size_t size_t;

size_t bank_balance = 300;

static std::mutex atm_mutex;
static std::mutex msg_mutex;
```

```

size_t withdraw(const std::string& msg, size_t amount) {
    std::lock(atm_mutex, msg_mutex);
    std::lock_guard<std::mutex> message_lock(msg_mutex, std::adopt_lock);

    std::cout << msg << " withdraws " << std::to_string(amount) << std::endl;

    std::lock_guard<std::mutex> account_lock(atm_mutex, std::adopt_lock);

    bank_balance -= amount;
    return bank_balance;
}

int main() {
    std::cout << "Starting balance is " << bank_balance << std::endl;

    std::future<size_t> bonnie = std::async(withdraw, "Bonnie", 100);
    std::future<size_t> clyde = std::async(deposit, "Clyde", 100);

    std::cout << "Balance after Clyde's deposit is " << clyde.get() << std::endl;
    std::cout << "Balance after Bonnie's withdrawal is " << bonnie.get() << std::endl;

    std::cout << "Final bank balance is " << bank_balance << std::endl;

    return 0;
}

```

Here we just show the withdraw function. In your homework downloads you will find the full source code in the file `bonnie_and_clyde.cpp`. Note that we have modified the code from lecture so that the deposit and withdrawal functions return a value, rather than just modifying the global state.

Compile and run the `bonnie_and_clyde` program several times. To compile use the command

```

c++ -std=c++11 -pthread bonnie_and_clyde.cpp -o bonnie_and_clyde.exe

```

Does it always return the same answer? On my laptop I ran the example and obtained the following four results (your machine may generate any combination of the following. Here are four example runs):

```

[0]$ ./bonnie_and_clyde.exe
Starting balance is 300
Balance after Clyde's deposit is Clyde deposits 100
400
Balance after Bonnie's withdrawal is Bonnie withdraws 100
300
Final bank balance is 300

[1]$ ./bonnie_and_clyde.exe
Starting balance is 300
Bonnie withdraws 100
Clyde deposits 100
Balance after Clyde's deposit is 300
Balance after Bonnie's withdrawal is 200
Final bank balance is 300

[2]$ ./bonnie_and_clyde.exe
Starting balance is 300
Balance after Clyde's deposit is Bonnie withdraws 100
Clyde deposits 100

```

```

300
Balance after Bonnie's withdrawal is 200
Final bank balance is 300

[3]$ ./bonnie_and_clyde.exe
Starting balance is 300
Clyde deposits Balance after Clyde's deposit is 100
Bonnie withdraws 100
400
Balance after Bonnie's withdrawal is 300
Final bank balance is 300

```

See if you can trace out how the different tasks are running and being interleaved. Notice that everything is protected from races – we get the correct answer at the end every time.

Note the ordering of calls in the program:

```

std::future<size_t> bonnie = std::async(withdraw, "Bonnie", 100);
std::future<size_t> clyde  = std::async(deposit, "Clyde", 100);

std::cout << "Balance after Clyde's deposit is " << clyde.get() << std::endl;
std::cout << "Balance after Bonnie's withdrawal is " << bonnie.get() << std::endl;

```

That is, the ordering of the calls is Bonnie makes a withdrawal, Clyde makes a deposit, we get the value back from Clyde’s deposit and then get the value back from Bonnie’s withdrawal. In run 0 above, that is how things are printed out: Bonnie makes a withdrawal, Clyde makes a deposit, we get the value after Clyde’s deposit, which is 300, and we get the value after Bonnie’s withdrawal, which is 200. Even though we are querying the future values with Clyde first and Bonnie second, the values that are returned depend on the order of the original task execution, not the order of the queries.

But let’s look at run 2. In this program, there are three threads: the main thread (which launches the tasks), the thread for the bonnie task and the thread for the clyde task. The first thing that is printed out after the tasks are launched is from Clyde’s deposit. But, before that statement can finish printing in its entirety, the main thread interrupts and starts printing the balance after Clyde’s deposit. The depositing function isn’t finished yet, but the `get()` has already been called on its future. That thread cannot continue until the clyde task completes and returns its value. The clyde task finishes its `cout` statement with “100”. But then the bonnie task starts running - and it runs to completion, printing “Bonnie withdraws 100” and it is done. But, the main thread was still waiting for its value to come back from the clyde task – and it does – so the main thread completes the line “Balance after Clyde’s deposit is” and prints “400”. Then the next line prints (“Balance after Bonnie’s withdrawal”) and the program completes.

Try to work out the sequence of interleavings for run 1 and run 3. There is no race condition here – there isn’t ever the case that the two tasks are running in their critical sections at the same time. However, they can be interrupted and another thread can run its own code – it is only the critical section code itself that can only have one thread executing it at a time. Also note that the values returned by bonnie and clyde can vary, depending on which order they are executed, and those values are returned by the `future`, regardless of which order the futures are evaluated.

std::future::get Modify your `bonnie_and_clyde` program in the following way. Add a line in the `main` function after the printout from Bonnie’s withdrawal to double check the value:

```

std::cout << "Balance after Bonnie's withdrawal is " << bonnie.get() << std::endl;
std::cout << "Double checking Clyde's deposit " << clyde.get() << std::endl;

```

(The second line above is the one you add.) What happens when you compile and run this example? Why do you think the designers of C++ and its standard library made futures behave this way?

3.2 Function Objects

One of the most powerful features of C++ is function overloading. As we have seen, one can essentially create new languages embedded within C++ by appropriately overloading functions and operators. Overloading `operator()` is particular interesting. Although we have overloaded it for matrix and vector types to be used as indexing, `operator()` can be defined to take any kind of arguments. In essence, overloaded `operator()` allows you to pass arbitrary parameters to a function that you define – in essence letting you build your own functions.

Now, of course, when you write a function down when you are programming, you are building functions. But when you build an object that has an overloaded `operator()`, your *program* can generate new functions. This is similar to the concept of lambda described below (and in fact was leveraged to develop the original Boost.Lambda library).

Objects that have an `operator()` defined are called *function objects*. What makes function objects so powerful is that they can be used just like a function in your programs. The syntax in C++ for a function call is a name followed by a parameter list enclosed in parentheses. This same syntax can mean a plain function call – or a call to an object's `operator()` member function.

We have seen two cases already where we pass functions as parameters: `std::thread` and `std::async`. But what is important in those cases is not that an actual function is being passed to `std::thread` and `std::async`, but rather that the thing being passed is *callable* (i.e., that it can be passed parameters and a function run on those parameters).

Let's look at an example. In the running Bonnie and Clyde bank deposit and withdrawal program, we have had two functions `withdraw` and `deposit` that are invoked asynchronously by the main program. Instead of defining `withdraw` and `deposit` as functions, let's instead define them as function objects. Let's start with a transaction class.

```
int bank_balance = 300;

class transaction {
public:
    transaction(int _sign, int& _balance) :
        sign(_sign), balance(_balance) {}

    void operator()(const std::string& msg, int amount) {
        std::lock_guard<std::mutex> tr_lock(tr_mutex);
        std::cout << msg << ": " << sign*amount;
        std::cout << " -- Balance: " << balance << std::endl;
        balance += sign*amount;
    }

    void operator()(const std::string& msg, int amount) {
        std::lock_guard<std::mutex> tr_lock(tr_mutex);
        balance += sign*amount;
    }

private:
    int sign; int &balance;
    static std::mutex tr_mutex;
};

std::mutex transaction::tr_mutex; // construct static member
```

The constructor for this class takes a sign (which should be plus or minus one and keeps a reference to a balance value. It also has its own mutex for safely modifying the balance.

Previously, the `withdraw` function was defined and invoked like this

```
void withdraw(int amount) {
    bank_balance -= amount;
```

```

}

int main() {

    withdraw(100);

    return 0;
}

```

We can create a `withdraw` function object that can be used exactly the same way:

```

int main() {
    transaction withdraw(-1, balance);

    withdraw(100);

    return 0;
}

```

The full Bonnie and Clyde example (with threads, say) would then be:

```

int bank_balance = 300;

int main() {
    transaction deposit ( 1, balance); // define an object named deposit
    transaction withdraw(-1, balance); // define an object named withdraw

    cout << "Starting balance is " << bank_balance << endl;

    thread bonnie(withdraw, "Bonnie", 100);
    thread clyde(deposit, "Clyde", 100);

    bonnie.join();
    clyde.join();

    cout << "Final bank balance is " << bank_balance << endl;

    return 0;
}

```

(For brevity we omit the messages that were printed by the original functions). This looks very much like the original program where `withdraw` and `deposit` were functions. But remember, in this case, they are objects. And they are created at run time when we construct them. In summary, function objects let you dynamically create and use stateful objects, just as if they were statically defined functions.

Read through the code for this example and run it (and/or modify it). You will find the full implementation in `bonnie_and_clyde_function_object.cpp` which can be compiled using the command

```

c++ -std=c++11 -pthread bonnie_and_clyde_function_object.cpp
    -o bonnie_and_clyde_function_object.exe

```

3.3 Lambda

Most of us are accustomed to programming in an imperative fashion. Programs are an ordered set of commands that govern what the CPU should process and in what order. Even in a parallel program written in an imperative fashion, each of the concurrent executions are an ordered set of instructions. In contrast, functional programming consists of expressions that are evaluated to produce a final value. In pure functional languages, there is no state per se, there is no notion of assignment. In functional programming, the

treatment of functions and data are unified. Expressions can evaluate to be functions, for example. That is, the value of an expression can be function, just as the value of an expression could also be a number or a string or a list. Using Scheme as an example, the following expression evaluates to be a function.

```
(lambda (x) (* x x))
```

The value of this expression is a function – a lambda expression. A function can be applied – in this case the function takes one parameter and evaluates the product of that parameter with itself (i.e., it squares it). Note that this function has no name – it just *is* a function. The function itself is distinct from the body of the function. The expression `(* x x)` is the value obtained by applying the multiplication operator to `x` and `x`.

An unnamed function in Scheme can be used just as a named function:

```
=> ((lambda (x) (* x x)) 7)
49
```

The expression above applies the function `(lambda x) (* x x)` to the value 7 and returns 49.

C++11 (and the Boost.Lambda library before it) provides support for lambdas – for unnamed functions – in C++. Besides the expressiveness in being able generate new functions from other functions, lambdas are convenient when one wants to pass a function to another function, such as when we use `std::async`.

In lecture we have been using matrix-vector product as a running example. The sequential function looks like this:

```
void matvec(const Matrix& A, const Vector& x, Vector& y) {
    double sum = 0.0;
    for (int i = 0; i < A.numRows(); ++i) {
        sum = 0.0;
        for (size_t j = 0; j < A.numCols(); ++j) {
            sum += A(i, j) * x(j);
        }
        y(i) = sum;
    }
}
```

We applied a simple parallelization to this nested loop by performing the inner loop as parallel tasks. However, to realize that we had to create a helper function and put the inner loop computation there.

```
double matvec_helper(const Matrix& A, const Vector& x, int i, double init) {
    for (size_t j = 0; j < A.numCols(); ++j) {
        init += A(i, j) * x(j);
    }
    return init;
}

void matvec(const Matrix& A, const Vector& x, Vector& y) {
    std::vector<std::future<double>> futs(A.numRows());
    for (int i = 0; i < A.numRows(); ++i) {
        futs[i] = std::async(matvec_helper, A, x, i, 0.0);
    }
    for (int i = 0; i < A.numRows(); ++i) {
        y(i) = futs[i].get();
    }
}
```

Unfortunately, now the logic of the original single algorithm is spread across two functions.

With a lambda on the other hand, we don't need to create a separate, named, helper function. We just need to create an anonymous function.

```

void matvec(const Matrix& A, const Vector& x, Vector& y) {
    std::vector<std::future<double>> futs(A.numRows());
    for (int i = 0; i < A.numRows(); ++i) {
        futs[i] = std::async( [&A, &x](int row) -> double {
            double sum = 0.0;
            for (size_t j = 0; j < A.numCols(); ++j) {
                sum += A(row, j) * x(j);
            }
            return sum;
        }, i);
    }
    for (int i = 0; i < A.numRows(); ++i) {
        y(i) += futs[i].get();
    }
}

```

In C++, lambdas begin with `[]`, which can optionally have “captured” variables passed (that is, variables outside of the scope of the lambda that will be accessed – in this case `A` and `x`). The value that the function is really parameterized on is the row we are performing the inner product with, so we make `row` a parameter to the lambda. The function returns a double, which we indicate with `->`. Finally we have the body of the function, which is just like the helper function body – and just like the inner loop we had before. Finally, we indicate that the argument to be passed to the function when it is invoked is `i`.

Since `i` is in the outer scope of the lambda, we can also capture it:

```

void matvec(const Matrix& A, const Vector& x, Vector& y) {
    std::vector<std::future<double>> futs(A.numRows());
    for (int i = 0; i < A.numRows(); ++i) {
        futs[i] = std::async( [&A, &x, i]() -> double {
            double sum = 0.0;
            for (size_t j = 0; j < A.numCols(); ++j) {
                sum += A(i, j) * x(j);
            }
            return sum;
        });
    }
    for (int i = 0; i < A.numRows(); ++i) {
        y(i) += futs[i].get();
    }
}

```

There are a variety of online references to learn more about C++ lambdas and I encourage you to learn more about – and to use – this powerful programming feature.

4 Exercises

The exercises in this assignment are centered on computing vector norms and matrix vector multiplication. Each of the graded code deliverables will ask you to sequentially fill in the following functions, located in the file `amath583.cpp`:

```

// == Functions FOR AMATH 483 & 583 =====
double twoNorm(const Vector& x);
double twoNormAscend(const Vector& x);
double twoNormDescend(const Vector& x);

double partitionedTwoNorm(const Vector& x, size_t partitions);

```



```

void ptn_worker(const Vector& x, std::size_t begin, size_t end, double& partial);

double recursiveTwoNorm(const Vector& x, size_t levels);
double rtn_worker(const Vector& x, std::size_t begin, std::size_t end, std::size_t level);

// == Functions For AMATH 583 Only =====
void task_matvec(const Matrix& A, const Vector& x, Vector& y, std::size_t partitions);
void matvec_helper(const Matrix& A, const Vector& x, Vector& y, std::size_t begin, std::size_t end

```

The python grader will be generate its own main file when it checks your implementations. To help you check for compilation errors without having to continually re-run the grader, we will provide you with the main files used by the grader. They are called

```

test_serial_norms.cpp // checks twoNorm, TwoNormAscend, and TwoNormDescend
test_parallel_norms.cpp // checks partionedTwoNorm, ptn_worker, recursiveTwoNorm, rtn_worker
test_parallel_matmul.cpp // tests task_matvec, matvec_helper

```

They should not be modified for this assignment. By running “make all” using the included makefile you can check to see whether your implementations compile successfully.

Note: the empty src files that you download for the assignment will compile and run but produce incorrect results since the implementations are empty. We recommend you check that your code compiles after modifying each function, rather than waiting till the end of the assignment.

Once you are confident that each of these three files compiles successfully, place all your code files in the candidate folder and run the autograder using the command

```
$ python2 grader.py
```

4.1 Three Serial Ways to Compute The 2-Norm

The file amath583.cpp contains three empty functions:

```

double twoNorm(const Vector& x){}
double twoNormAscend(const Vector& x){}
double twoNormDescend(const Vector& x){}

```

Each of these functions compute the Euclidean 2-norm of the vector x in the following ways:

1. `twoNorm(x)` computes the 2-norm of x in the order it natural ordering. (i.e. $\|x\|_2 = \sqrt{\sum_{i=1}^n |x_i|^2}$)
2. `twoNormAscend(x)` computes the 2-norm of the vector x with the values in ascending order (i.e. smallest to largest). You may find the function `std::sort` to be a useful.
3. `twoNormDescend(x)` computes the 2-norm of the vector x with the values in descending order (i.e. largest to smallest).

Code Deliverable

Complete the implementations for `twoNorm`, `twoNormAscend`, and `twoNormDescend` in the file `amath583.cpp`.

Written Deliverable

Create a program `prob4_1.exe` to test out each of your norm implementations on random vectors of dimension 1000. You should use the function `randomVector(n)` provided in `amath583.cpp` to generate the vectors. Do each of the implementation produce exactly the same result if tested on the same vector? Make sure you check all 15 digits of your computed norms. In your written response pdf, explain the results you are seeing in at most three sentences. We will **not** run `prob4_1.exe` using the autograder so you can write it however you like.

Parallelizing the Norm

In the next two exercises you will be adding concurrency and (hopefully) parallelism to speed up computation of the norm. This will be accomplished by writing a partitioned norm and a recursive norm.

TIP: It might be helpful to first write a serial version of partitioned norm and recursive norm. Once these are compiling and running, then you can add parallelism.

4.2 Norm (Partitioned Version)

Inside the file `amath583.cpp` you will find the following empty functions

```
double partitionedTwoNorm(const Vector& x, size_t partitions);
void ptn_worker(const Vector& x, size_t begin, size_t end, double& partial);
```

`partitionedTwoNorm` takes the vector `x` as an argument, computes its two norm (Euclidean norm) and returns that value. However, instead of computing the 2-norm by summing over all elements of the vector directly, the full sum is partitioned into smaller sums. For example, for two partitions this would be

$$\underbrace{\sum_{j=0}^{n-1} x_i^2}_{\text{full sum}} = \underbrace{\sum_{j=1}^m x_i^2}_{\text{partition 1}} + \underbrace{\sum_{j=m+1}^{n-1} x_i^2}_{\text{partition 2}} \quad \text{for} \quad m = \lfloor \frac{n-1}{2} \rfloor$$

For your implementation, you are to use `std::thread` to execute worker tasks to compute each partial sum. The number of worker tasks (the number of partitions to make) is given by the argument `partitions`. Each worker task should have the prototype:

```
void ptn_worker(const Vector& x, size_t begin, size_t end, double& partial);
```

The function should *safely* (i.e., no race conditions) accumulate the sum of squares of the elements in `x` between `begin` and `end` (i.e. for indices i that satisfy $\text{begin} \leq i < \text{end}$) and add this sum to the value `partial`.

You may use synchronization to safely do the accumulation into a single variable or you may accumulate into separate variables that are later summed together. It is important that the workers just compute and accumulate sums of squares and then the sum of all of those sums is computed – and then the square root taken on that final value. The `pi` example we worked through in lecture should be helpful guidance for this exercise.

Note that there are several ways to choose partitions. Ideally, the number of partitions evenly divides into the length of the vector. When this is not the case, it is possible that certain partitions will have more elements than others. For example, suppose that we have n total tasks (indexed by $0, \dots, n-1$) and p processors (index by $0, \dots, p-1$); one of the simplest ways to choose the tasks (though not the most efficient) is if the j th processor completes the tasks

$$\begin{cases} j\Delta \text{ through } (j+1)\Delta - 1 & \text{if } j < p-1 \\ j\Delta \text{ through } n-1 & \text{if } j = p-1 \end{cases} \quad \text{where} \quad \Delta = \lfloor n/p \rfloor.$$

For this assignment, `partitionedTwoNorm` should divide vector indices in this way. If we have p total partitions (indexed by $0, \dots, p-1$) and a vector of dimension n then the j th partition should sum the squares of the indices

$$\begin{cases} [j\Delta, (j+1)\Delta) & j < p \\ [j\Delta, n) & j = p \end{cases} \quad \text{where} \quad \Delta = \lfloor n/p \rfloor.$$

Code Deliverable

Complete the implementations for `partitionedTwoNorm` and `ptn_worker` in the file `amath583.cpp`.

4.3 Norm (Recursive Version)

Another approach (common in the functional programming community) for decomposing a problem into smaller parts is *recursion*. Recursion is an interesting issue in functional programs because the programs are simply functions. Whereas it isn't too difficult in the imperative world to think of a function calling itself, in the functional world, recursion means that a function is *defined* in terms of itself. For example, we could define a function of $f(x)$ in the following way

$$f(x) = x \times f(x - 1)$$

So what is f ? It's defined in terms of itself. And, what do we get when we plug in a value for x , say, 3?

$$f(3) = 3 \times f(2) = 3 \times 2 \times f(1) = 3 \times 2 \times 1 \times (0) = \dots$$

To make a self-referential definition sensible we have to define a specific value for the function with some specific argument. This is usually called the "base case". In this example, we could define $f(0) = 1$. Then the function is defined as:

$$f(x) = \begin{cases} 1 & \text{if } x = 0 \\ x \times f(x - 1) & \text{otherwise} \end{cases}$$

One interesting way to think about recursion is to think of it in terms of abstraction. In the above example, we don't define what $f(x)$ is. Rather we define it as x times $f(x - 1)$ – that is, we are abstracting away the computation – we are pretending that we know the value of $f(x - 1)$ when we compute the value of $f(x)$.

This interpretation is also useful for decomposing a problem into smaller pieces that we assume we can compute. For instance, suppose we want to take the sum of squares of the elements of a vector (as a precursor for computing the Euclidean norm, say). Well, we can abstract that away by pretending we can compute the sum of squares of the first half of the vector and of the second half of the vector. The sum of squares of the entire vector is the sum of the values from the two halves. For example:

```
double sum_of_squares(const Vector& x, size_t begin, size_t end) {
    return sum_of_squares(x, begin,                begin+(end-begin)/2)
        + sum_of_squares(x, begin+(end-begin)/2, end) ;
}
```

To use an over-used phrase – see what we did there? We've defined the function as the result of that same function on two smaller pieces. This general approach, particularly when dividing a problem in half to solve it, is also known as "divide and conquer."

Now, in this code example, as with the mathematical example, we need some non-self-referential value. In this case there are several options. We could stop when the difference between `begin` and `end` is sufficiently small. Or, we can stop after we have descended a specified number of levels. In the latter case:

```
double sum_of_squares(const Vector& x, size_t begin, size_t end, size_t level) {
    if (level == 0) {
        return // code that computes the sum of squares of x from begin to end
    } else
        return sum_of_squares(x, begin,                begin+(end-begin)/2, level-1)
            + sum_of_squares(x, begin+(end-begin)/2, end,                , level-1) ;
}
```

Now, we can parallelize this divide and conquer approach by launching asynchronous tasks for the recursive calls and getting the values from the futures. Consider for a second what will happen though particularly if we use `std::launch::deferred`. We won't actually do any computation until the entire tree of tasks is built. The first call will create two futures, when `get` is called on them, each of them will create two more futures that don't compute, and so on until the base case is reached – at which point the tree collapses back upwards.

Inside the file `amath583.cpp` you will find the following empty functions

```
double recursiveTwoNorm(const Vector& x, size_t levels);
double rtn_worker(const Vector& x, size_t begin, size_t end, size_t level);
```

The function `recursiveTwoNorm` takes `x` as an argument, computes its two norm (Euclidean norm) and returns that value. The number of levels to recurse is given by the argument `levels`. At each level of recursion, your function should divide your vector into roughly half. It may be easier (but isn't necessary) to implement your recursion function first in serial before adding in parallelization. For your final implementation, you are to use `std::async` to execute worker tasks. Each worker task should have this prototype:

```
double rtn_worker(const Vector& x, size_t begin, size_t end, size_t level);
```

The function should *safely* accumulate the sum of squares of the elements in `x` between `begin` and `end` (the half-open interval) by asynchronously invoking itself with intervals of half the size (as shown above).

Code Deliverable

Complete the implementations for `recursiveTwoNorm` and `rtn_worker` in the file `amath583.cpp`.

4.4 Analyzing the Performance of Recursive Norm and Partitioned Norm

In this section you will be asked to generate several performance plots. You may use any software of your choice (e.g. Excel, Matlab, Google Sheets). We will **not** use our autograder to test the programs you write for generating the data, so you can write each file however you like. For each timing data point you should average over at least 100 trials. You may generate all random vectors using the function `randomVector` provided in `amath583.cpp`. For timing your code you can use the class located inside `HRTimer.hpp`.

Written Deliverable

1. Write the program `prob4_4a.exe` that computes the norms of randomly generated vectors of size

$$2, 4, 8, \dots, 2^{23}$$

using `twoNorm()`, `partitionTwoNorm()` with 4 partitions, and `recursiveTwoNorm()` with 4 levels. Time each of these cases and create a plot of running time vs vector size for each function. Label and include all three plots in your solution pdf. Explain your results in less than two sentences.

2. Write the program `prob4_4b.exe` main program that computes the norms of random vectors of size 2^{15} using `partitionedTwoNorm()` with 1, 3, 5, ..., 19 partitions. Time your results and produce a plot in your solution pdf showing running time vs number of partitions. Explain your results in less than two sentences.
3. Write the program `prob4_4c.exe` that computes the norms of random vectors of size 2^{15} using `recursiveTwoNorm()` with 1, 2, 3, ..., 10 levels. Time your results and produce a plot in your solution pdf showing running time vs number of levels. Explain your results in less than two sentences.

4.5 Matrix Vector Product (AMATH 583 only)

In lecture we discussed some approaches for parallelizing dense matrix-vector products. Inside the file `amath583.cpp` you will find the following empty functions

```
void task_matvec(const Matrix& A, const Vector& x, Vector& y, size_t partitions);
void matvec_helper(const Matrix& A, const Vector& x, Vector& y, size_t begin, size_t end);
```

The function `task_matvec` multiplies the matrix `A` and the vector `x` and writes the result to the vector `y`. The function should divide the overall problem into the given number of partitions (using threads or tasks). Your helper function `matvec_helper` should be prototyped as described above.

Code Deliverable

Complete the implementations for `task_matvec` and `matvec_helper` in the file `amath583.cpp`.

Written Deliverable

Write the program `prob4_5.exe` to time matrix vector products for square matrices and vectors of size

$$2, 4, \dots, 2^{13}.$$

In your solution pdf, add a single plot showing timing results for: (1) the non-parallelized `matvec` that is implemented in `amath583.cpp` as `matvec`, (2) `task_matvec()` with 2 partitions, and (3) `task_matvec()` with 4 partitions. For each time point on the plot, you should average over 100 trials. We will **not** use our autograder to test `prob4_5.exe`, so you can write this file however you like.

Extra Credit. Implement task-based matrix-vector product using C++ lambda rather than the helper function. You should put the code in `amath583.cpp`. Create a program `prob4_5EC.exe` to time your code and produce an additional graph in your writeup that is made the same way as described in the previous paragraph.

4.6 Additional Written Questions

1. You measured bandwidth explicitly in the bandwidth program and implicitly with the roofline program. Explain in no more than one paragraph how well do the bandwidths from these two profiles match? Include the two generated files (one pdf and png do not rename them unless the extensions are different) with your other materials for this assignment.
2. **Extra Credit.** Calculate the numeric intensity of sparse matrix-vector product with compressed sparse row (CSR) format, for the two cases of index types being integer (32 bit) and `size_t` (64 bit). Use your roofline graph to estimate the performance of CSR sparse matrix vector product on your computer.
3. **Extra Credit.** Implement sparse matrix-vector product with CSR format and compare the results you obtain experimentally with what was predicted by your roofline model.

5 Turning in The Exercises

1. All your written responses should be uploaded to canvas as a pdf file to the written assignment. If necessary you should include a list of references (electronic, written, human).
2. All your `.cpp`, `.hpp` files, and `Makefile`, should go in the tarball `ps5.tgz` and uploaded to the coding assignment. Do not tar the candidate folder, or the entire grader! Only tar your code files. Your command for tarring files should look something like this

```
tar -czf ps5.tgz *.cpp *.hpp Makefile
```

AMATH 483: Before submitting homework, check that the command “make all” generates the following executables:

1. `test_serial_norms.exe`
2. `test_parallel_norms.exe`
3. `prob4_1.exe`
4. `prob4_4a.exe`, `prob4_4b.exe`, `prob4_4c.exe`

Note: The first two files were provided for you in the homework and should not be modified. If you find that they no longer compile, then the problem is due to changes you made in `amath583.cpp`.

AMATH 583: Before submitting homework, check that the command “make all” generates the following executables:

1. `test_serial_norms.exe`
2. `test_parallel_norms.exe`
3. `test_parallel_matmul.exe`
4. `prob4_1.exe`
5. `prob4_4a.exe`, `prob4_4b.exe`, `prob4_4c.exe`
6. `prob4_5.exe`

Note: The first three files were provided for you in the homework and should not be modified. If you find that they no longer compile, then the problem is due to changes you made in `amath583.cpp`.

6 Learning Outcomes

At the conclusion of week 5 students will be able to

1. Characterize the differences between a process and a thread
2. Use `std::thread` to spawn a thread
3. Describe the arguments passed to `std::thread`
4. Describe what `join()` does
5. Write a simple program that launches four concurrent threads and that demonstrate concurrency in some fashion
6. Explain the difference between concurrency and parallelism
7. Define data race and describe a scenario where a data race would occur
8. Write a simple program that exhibits a data race
9. Name two hardware supported atomic operations and show how to use them to protect a critical region
10. Describe a spin lock and its advantages and disadvantages
11. Identify potential data races in a concurrent program
12. Correctly use `std::mutex()` to protect critical regions
13. Explain the difference between `std::mutex()` and a spin lock
14. Define RAII and why it is used
15. Explain `std::lock_guard`
16. Correctly use `std::lock_guard` to protect a critical section
17. Explain what deadlock is
18. Describe a scenario where deadlock might occur

19. Use `std::thread_guard` and `std::lock` to correctly protect a critical section having multiple mutexes – in a deadlock-free fashion

At the conclusion of week 6 students will be able to

1. Describe the four steps of creating a parallel program (as presented by Mattson et al).
2. Describe the difference between `std::thread` and `std::async`
3. Describe the two different modes for launching `std::async` and what the difference is between them
4. Write a simple program that launches four concurrent tasks with `std::async` and that demonstrate concurrency in some fashion
5. Describe what `std::future` is
6. Define atomicity as it relates to concurrent programs
7. Use `std::atomic` to protect a shared variable
8. Characterize the types that can be used with `std::atomic`
9. Describe a parallelization approach for parallelizing inner product, dense matrix-vector product, sparse matrix-vector product, and dense matrix-matrix product
10. Describe the roofline performance model and an experimental approach that you could use to obtain its parameters
11. Use the roofline performance model to estimate the performance of dense matrix-vector product and sparse matrix-vector product with COO and CSR matrix formats